

**FORECASTING COVID-19 WITH
TEMPORAL HIERARCHIES AND ENSEMBLE METHODS**

A Thesis Presented

by

LI SHANDROSS

Submitted to the Graduate School of the
University of Massachusetts Amherst in partial fulfillment
of the requirements for the degree of

MASTER OF SCIENCE

May 2023

Biostatistics and Epidemiology

FORECASTING COVID-19 WITH TEMPORAL HIERARCHIES AND ENSEMBLE METHODS

A Thesis Presented

by

LI SHANDROSS

Approved as to style and content by:

Nicholas G. Reich, Chair

Evan L. Ray, Member

Benjamin W. Rogers, Member

Lisa Chasan-Taber, Department Chair
Biostatistics and Epidemiology

ABSTRACT

FORECASTING COVID-19 WITH TEMPORAL HIERARCHIES AND ENSEMBLE METHODS

MAY 2023

LI SHANDROSS

B.Sc., UNIVERSITY OF MASSACHUSETTS AMHERST

M.S., UNIVERSITY OF MASSACHUSETTS AMHERST

Directed by: Professor Nicholas G. Reich

Infectious disease forecasting efforts underwent rapid growth during the COVID-19 pandemic, providing guidance for pandemic response and about potential future trends. Yet despite their importance, short-term forecasting models often struggled to produce accurate real-time predictions of a complex and rapidly changing system. This gap in accuracy, which persists three years into the pandemic, warrants the exploration and testing of new methods to glean fresh insights.

In this work, we examine the application of the temporal hierarchical forecasting (THieF) methodology to probabilistic forecasts of COVID-19 incident hospital admissions in the United States. THieF is a recent, innovative forecasting technique that aggregates time-series data into a hierarchy made up of different temporal scales, produces forecasts at each level of the hierarchy, then reconciles those forecasts using optimized weighted forecast combination. While THieF's unique approach has shown substantial accuracy improvements in a diverse range of applications, such as opera-

tions management and emergency room admission predictions, this technique has not previously been applied to outbreak forecasting.

We generate candidate models formulated using the THieF methodology, which differ by their hierarchy schemes and data transformations, and weighted ensembles of the THieF models. The models are validated using weighted interval score (WIS), and the top-performing subset is compared to several benchmark models. These models include simple SARIMA models, a naive baseline model, and a median ensemble of operational incident hospitalization models from the US COVID-19 Forecast Hub. The THieF models and THieF ensembles demonstrate improvements in WIS and MAE, as well as competitive prediction interval coverage, over the benchmark models for both the validation and testing phases. These accuracy improvements suggest the THieF methodology may serve as a promising new technique for infectious disease forecasting that makes better predictions in real-time.

TABLE OF CONTENTS

	Page
ABSTRACT	iii
LIST OF TABLES	vii
LIST OF FIGURES	ix
 CHAPTER	
1. INTRODUCTION	1
1.1 Infectious disease forecasting	1
1.2 COVID-19 forecasting	1
1.2.1 US COVID-19 Forecast Hub	1
1.3 Temporal hierarchical forecasting	1
1.4 Goals	1
2. METHODS	2
2.1 Surveillance data	2
2.2 Forecasts	3
2.2.1 Forecast horizons	4
2.2.2 Forecast locations	4
2.2.3 Forecast quantiles	4
2.3 Date range of analysis	5
2.4 Benchmark comparison models	6
2.5 Temporal hierarchical forecasting (THieF)	8
2.6 Model specifications	8
2.6.1 THieF models	8
2.6.2 THieF ensemble models	8
2.6.3 ARIMA and SARIMA models	8

2.7	Metrics and Evaluation	8
2.7.1	Validation phase	9
2.7.2	Testing phase	9
3.	RESULTS	10
3.1	Validation phase.....	10
3.2	Testing phase	10
3.2.1	Summary of comparison models	10
3.2.2	Overall model performance	10
3.2.3	Model performance by horizon week	15
3.2.4	Model performance by pandemic wave	15
3.2.5	Model performance by location.....	15
3.2.6	Model performance by forecast week	15
4.	DISCUSSION	18
 APPENDICES		
	APPENDIX: SUPPLEMENTAL MATERIALS	19
	 BIBLIOGRAPHY	 24

LIST OF TABLES

Table		Page
2.1	Table of the 23 quantiles required by the Forecast Hub and their associated prediction interval bound.	5
3.1	Summary of overall model performance for the US national level, aggregated across forecast week and 1 to 4 week ahead horizons. Prediction interval (PI) coverage rates measure the fraction of times the fraction of times the 50% and 95% PIs covered the true values. A well-calibrated model should achieve close to nominal (0.5 and 0.95, respectively) coverage rates. The “rwis” and “rmae” show relative weighted interval score and relative mean absolute error, which compare each model to the baseline. The baseline is defined to have relative metrics of 1. Models that perform better than the baseline have relative WIS or MAE less than 1. The table is ordered by ascending WIS value. Thus two THieF-4wk models show the best performance out of the THieF models, meeting the baseline for both relative WIS and MAE.	10
3.2	Summary of overall model performance averaged over the states level. A similar ordering of models is seen, with the THieF-4wk-4root model performing the best and the THieF-12wk models performing the worst. However, the THieF-8wk models perform better in terms of model ranking for the states level. Additionally, relative WIS sees smaller values and 50% and 95% PI coverage see values closer to the nominal levels compared to that at the US national level, although relative MAE is larger with every model performing worse than the baseline.	15
1	Summary of model performance stratified by horizon week for the US national level, aggregated across forecast week. The results from this table were used to create Figure 4 that summarized average WIS by horizon week for the US national level and the states level. The table is ordered by ascending horizon week and WIS, with THieF-4wk-4root having the lowest WIS except for the 4 week ahead horizon. The baseline consistently does not have the best WIS.	20

2	Summary of model performance stratified by horizon week for the states level, aggregated across forecast week. The results from this table were used to create Figure 4 that summarized average WIS by horizon week for the US national level and the states level. The table is ordered by ascending horizon week and WIS, with THieF-4wk-4root having the lowest WIS except for the 3 week (although it is second lowest) and 4 week ahead horizon. The baseline again consistently does not have the best WIS, performing progressively worse compared to the other models for each additional week ahead horizon.	21
4	Summary of model performance stratified by horizon week and pandemic waves for the US national level, aggregated across forecast week. The results from this table were used to create Figure 5 that summarized average WIS by horizon week and pandemic wave for the US national level and the states level. The table is ordered by ascending pandemic wave, horizon week, and WIS, with THieF-4wk-4root having the lowest WIS for all horizons of the alpha wave and shorter horizons of the Winter 21 wave. The THieF-12wk-noTrans and THieF-12wk-4root models have the lowest and second lowest WIS for 3 week ahead and 4 week ahead horizons during the Winter 21 wave. The baseline consistently does not have the best WIS.....	22
6	Summary of model performance stratified by horizon week and pandemic waves for the states level, aggregated across forecast week. The results from this table were used to create Figure 5 that summarized average WIS by horizon week and pandemic wave for the US national level and the states level. The table is ordered by ascending pandemic wave, horizon week, and WIS, with THieF-4wk-4root again having the lowest WIS for all horizons of the alpha wave and shorter horizons of the Winter 21 wave. The THieF-12wk-noTrans has the lowest WIS for 3 week ahead and 4 week ahead horizons during the Winter 21 wave. The baseline only has the best WIS for the 1 week ahead horizon during the Winter 21 wave.	23

LIST OF FIGURES

Figure		Page
2.1	Overview of the period of analysis included in the paper, superimposed over COVID-19 incident hospitalizations for the entire United States. Vertical dashed lines indicate divisions between phases. The training phase begins on Monday, July 27, 2020 and ends on Monday, July 5, 2021. This phase covers 343 days, or 49 weeks. Forecasts for the training phase are made starting on Monday, December 7, 2020. The testing phase, although not included in this paper, begins Tuesday, July 6, 2021 and ends Monday, November 1, 2021. The division is based on including two waves (the Winter 2020-21 and the Alpha variant) of the pandemic in the training phase and one (Delta variant) in the testing phase. July 27, 2020 marks the beginning of when the forecast locations of interest had reliable truth data.....	7
3.1		11
3.2	Average WIS by the horizon week for each model for the US national level and the states level. The figures show similar trends in WIS when forecasts are stratified by horizon week compared to being aggregated across horizon weeks. All models show worse performance relative to the baseline at the 1-week ahead horizon, likely a result of being overconfident at this horizon. Once again, the THieF-4wk-4root model performs the best for WIS. Plots for average MAE are not shown due to their similarity to the average WIS plots.....	12

3.3	Average WIS by the horizon week and pandemic wave for each model for the US national level and the states level. The figures show similar trends in WIS when forecasts are stratified by horizon week and pandemic wave compared to being stratified only across horizon week. Again, all models show worse performance relative to the baseline at the 1-week ahead horizon likely due to being overconfident at this horizon. Model performance is more similar among the top four models during the Alpha wave compared to during the Winter21 wave. Notably, although the THieF-4wk-4root models largely show the lowest WIS, longer horizons during the Winter21 wave saw better performance by the poor-performing THieF-12wk models. Plots for average MAE are not shown due to their similarity to the average WIS plots.	13
3.4	Relative WIS by location for each model across all horizons for the entire testing phase. Each square contains the average relative WIS or a model at a particular location, colored based on the relative WIS compare to the baseline. Blue indicates model outperformed the baseline for that location on average, while red indicates the model performed worse than the baseline. Locations are sorted based on cumulative hospitalizations during the testing phase while models are sorted from lowest to highest overall relative WIS, aggregated across states (Table 2). The top three models by relative WIS overall during the testing phase outperformed the baseline in every location.	14
3.5	Average h-week ahead WIS and 95% PI coverage for each model for the US national level. Model performance differs from week to week: the WIS is the lowest during weeks with stable and low incident hospitalization values. The THieF-4wk models show the best performance with lower or comparable WIS to the baseline during the majority of the training period with especially good performance for 1-week ahead horizons. Meanwhile, the THieF-12wk showed lower or comparable WIS to the baseline during high incidence and changed points, which have been historically found difficult to forecast accurately. Plots for h-week ahead MAE are not shown due to their similarity to the WIS plots. All THieF models show very variable coverage rates, especially for a 1 week ahead horizon, only the models without a data transformation showing stable coverage starting around February 2021.	16

3.6	Average h-week ahead WIS and 95% PI coverage for each model for the states level. Once again, model performance differs from week to week with the lowest WIS values during weeks with stable and low incident hospitalization values and the THieF-4wk models showing the best performance. However, the THieF-12wk did not show good performance during high incidence and changed points for the states level; instead, THieF-12-4root actually showed some strangely large WIS values during April and May 2021. Plots for h-week ahead MAE are not shown due to their similarity to the WIS plots. One-week ahead coverage rates generally improved as time went on, the best models gaining stable and fairly good coverage rates starting in February 2021. For the 4-week ahead horizon, models without a data transformation showed the best performance and stable coverage starting around February 2021, though THieF-4wk-4root and THieF-8wk-4root also saw decent and fairly stable coverage rates after the downward Winter 21 change point.	17
-----	--	----

CHAPTER 1

INTRODUCTION

- 1.1 Infectious disease forecasting**
- 1.2 COVID-19 forecasting**
 - 1.2.1 US COVID-19 Forecast Hub**
- 1.3 Temporal hierarchical forecasting**
- 1.4 Goals**

CHAPTER 2

METHODS

2.1 Surveillance data

This analysis uses confirmed hospital admissions from HealthData.gov from the COVID-19 Reported Patient Impact and Hospital Capacity by State Time Series. This dataset reports state-aggregated daily incident hospitalizations starting on January 1, 2020 for U.S. states and territories. The data come from three primary sources: HHSTeleTracking, state/territorial health departments which report to HHS Protect, and National Health Care Safety Network (prior to July 15, 2020) [7].

On December 1, 2020, the U.S. COVID-19 Forecast Hub designated this dataset as the official ground truth data for COVID-19 incident hospitalizations. (No other sources were designated official ground truth data prior to this date.) [2] This meant that the HealthData.gov data was used to evaluate the performance of real-time incident hospitalization forecasts submitted to the Forecast Hub. Hence, teams were encouraged to use this hospitalization data set to train their models [3]. In the analysis, we also use this dataset as both input training data for our models and truth data against which to perform accuracy evaluations.

Hospitalization reporting was sporadic for the first four months of the pandemic, with consistent, reliable values for major states and territories achieved only in late July 2020. Specifically, reports of daily confirmed hospitalizations for all 50 states, Washington D.C., Puerto Rico, and the U.S. as a whole began on July 27, 2020. Although COVID-19 hospitalization reporting stabilized after that for cases and deaths, hospitalization values are often more reliable due to requirements for Medicare reim-

bursement. COVID-19 cases have been vastly underreported for the entire duration of the pandemic [5], and this problem has worsened with dwindling testing. Further, truth data collection and reporting for COVID-19 incident cases and deaths has been discontinued by Johns Hopkins University (the official ground truth data source) as of March 10, 2023 [4].

Surveillance data, like the HealthData.gov confirmed hospitalizations, sometimes required revisions due to reporting anomalies. Hospitalizations might be entered on incorrect dates or added after COVID-19 was determined to be the cause of hospitalization. Updates to values usually occurred within the first week or so of posting. As a result, real-time forecast models might train on data that include incident hospitalization numbers for location-date combinations that do not reflect the finalized values.

We query these confirmed hospitalizations from the ‘covidData’ package through the ‘covidHubUtils’ package, both developed by Reich Lab at the University of Massachusetts Amherst to deal with COVID-19 surveillance data. Versions of this data are stored such that users may access the truth data as of a specified date. For example, it is possible to retrieve of version of the confirmed hospitalization data set as of July 27, 2020. To ensure fair comparisons between the models developed for this analysis and some real-time models, our models are trained on the confirmed hospitalizations that would have been available as of the “forecast date” of the retrospective forecast.

2.2 Forecasts

A forecast for a target-location combination (e.g. “1 day ahead incident hospitalizations in Massachusetts”) is defined by one or both of a point forecast and a probabilistic forecast, represented by a set of 23 quantiles [3]. More information is

given about forecast targets, horizons, locations, and quantiles in the following subsections.

2.2.1 Forecast horizons

Short-term forecasts for COVID-19 have been defined to encapsulate horizons of 1- to 4-weeks ahead, or, equivalently, 1- to 28-days ahead [3]. Forecasts for incident hospitalizations are made on a daily time scale, so we choose to restrict¹ the forecast horizons we evaluate to 1- to 28-days ahead.²

2.2.2 Forecast locations

We make forecasts for 53 of 55 locations from the HealthData.gov confirmed hospitalizations dataset during the period of analysis. These locations include the 50 states, four jurisdictions/territories (Washington D.C., Puerto Rico, American Samoa, and Guam), and a U.S. national level. American Samoa and Guam are excluded due to their very low or zero COVID-19 incident hospitalizations over the entire pandemic. Forecasts for low-count locations that predicted single-digit or zero hospitalizations were often very accurate; hence, these locations do not offer meaningful contributions to our understanding of differences in forecast skill between modeling approaches [3].

2.2.3 Forecast quantiles

The Forecast Hub specifies the usage of 23 quantiles

$$\mathbb{Q} = \{.010, 0.025, .050, .100, \dots, .900, .950, .990\}$$

to represent a full probabilistic distribution, with the median 0.500 quantile usually taken as the model’s point forecast. We follow this construction, though with a slight

¹Some models make up to 84-day ahead forecasts due to the THieF methodology construction.

²Note, there is a slight exception to this restriction. See Metrics and Evaluation for more information.

modification: prediction intervals are used to define a full distribution of forecasts. The bounds of these intervals can be used to obtain the desired quantiles, as shown in the table below. For example, the values for a 0.500 and a 0.950 quantile (see the Quantile row) are obtained by taking the lower and upper bounds of the 95% prediction interval (see the Interval row, where "L" and "U" represent the lower and upper bounds, respectively).

Quantile	.010	.025	.050	0.100	0.150	...	0.500	...	0.900	0.950	.975	.990
Interval	98L	95L	90L	80L	70L	...	Median	...	80U	90U	95U	98U

Table 2.1. Table of the 23 quantiles required by the Forecast Hub and their associated prediction interval bound.

2.3 Date range of analysis

We focused on the period between Monday, July 27, 2020 and Sunday, October 2, 2022. This date range spans 798 days, or 114 weeks, which is just over 26 months. The entire period of analysis is split into two main phases: a validation phase of 462 days (66 weeks) and a testing phase of 236 days (48 weeks). The split between these two phases is such that the validation phase ends on Sunday, October 31, 2021 and the testing phase begins on Monday, November 1, 2021.

The first day of the period of analysis (July 27, 2020) was selected as the first day of reliable truth data for all locations of interests. Forecasting for the validation phase began on Monday, December 7, 2020, as this is the first Monday after the Health-Data.gov confirmed hospitalizations data set was declared the official ground truth data. The division between the validation and testing phases was chosen to encapsulate three COVID-19 variant-based waves in the validation phase (Winter 2020-21, Alpha, and Delta) and two waves in the testing phase (Omicron and BA.4/BA.5).

The separation between pandemic waves is made based on U.S. national level incident hospitalizations, pictured below in [Figure X].

2.4 Benchmark comparison models

Two models from the Forecast Hub—the COVIDhub-baseline and COVIDhub-4_week_ensemble—are selected to serve as comparisons for the models in this analysis during the testing phase evaluation. Many of the metrics used in this evaluation are not especially meaningful on their own, requiring a relative comparison among models forecasting the same targets (1 to 28 day ahead incident hospitalizations for COVID-19).

The COVIDhub-baseline: A naive baseline model created and run by the US COVID-19 Forecast Hub meant to serve as a benchmark model to compare against. This model makes forecasts using the previous day’s incident hospitalization value as the point forecast (and 0.5 quantile for the probabilistic forecast) and past changes in daily incidence to calibrate uncertainty in its probabilistic forecasts. Beating the baseline is typically seen as one marker of a useful model [3]

The COVIDhub-4_week_ensemble: A median ensemble of all the submitting models for incident hospitalizations at the COVID-19 Forecast Hub, run by the Hub team that includes forecasts up to 28 day (4 week) ahead horizons. This model produces the same 1 to 14 day (1 to 2 week) ahead horizon forecasts as the COVIDhub-ensemble model, but also forecasts longer horizons. These longer horizons were removed from the main ensemble model due to research showing a sharp decline in accuracy in forecasts predicting beyond 2 week aheads (see Forecast horizons). However, these longer horizons are useful as a comparison for research purposes. Note that this ensemble model included slightly different contributing models based on what models were submitting to the COVID-19 Forecast Hub at a different times (check for specifics). [1], [6]




Figure 2.1. Overview of the period of analysis included in the paper, superimposed over COVID-19 incident hospitalizations for the entire United States. Vertical dashed lines indicate divisions between phases. The training phase begins on Monday, July 27, 2020 and ends on Monday, July 5, 2021. This phase covers 343 days, or 49 weeks. Forecasts for the training phase are made starting on Monday, December 7, 2020. The testing phase, although not included in this paper, begins Tuesday, July 6, 2021 and ends Monday, November 1, 2021. The division is based on including two waves (the Winter 2020-21 and the Alpha variant) of the pandemic in the training phase and one (Delta variant) in the testing phase. July 27, 2020 marks the beginning of when the forecast locations of interest had reliable truth data.

The forecasts for the comparison models are obtained from the Zoltar forecast archive, a research data repository developed by Reich Lab.

2.5 Temporal hierarchical forecasting (THieF)

- Add more mathematical detail here - Aggregation: We start with data at a daily scale. Then, based on the aggregation levels of interests for a particular model (e.g. weekly, every two weeks, every four weeks, etc), we induce this sort of artificial aggregation so that we have vectors of data add each of these time scales. Next, forecasts are made using each of these datasets separately. The reason that we are interested in different time scales is that there may be important patterns that are lost add higher time scales amid the noise of lower ones. More often than not, different patterns will be picked up at different aggregation levels, which means we get forecasts for the same future time periods that may not agree. This is where reconciliation comes into play, in which we combine all of the forecasts that we made.

2.6 Model specifications

2.6.1 THieF models

2.6.2 THieF ensemble models

2.6.3 ARIMA and SARIMA models

2.7 Metrics and Evaluation

Despite being made on a daily scale, hospitalization forecasts have been evaluated on a weekly scale by the Forecast Hub, and we follow this convention. However, modeling teams have made forecasts on different days (e.g. Saturday vs Sunday vs Monday), leading to a mismatch between what day an h -day ahead forecast is predicting for. Hence, for the models we construct, we actually use 3- to 30-day ahead horizons to ensure all forecasts are predicting for the same dates. It should

be noted, however, this may result in a slight disadvantage for models whose 3- to 30-day ahead forecasts are being compared to 1- to 28-day ahead forecasts.

2.7.1 Validation phase

- Serves as time during which to train the candidate models, calibrate the forecasts, and to select the best models to pass on to the testing phase - All validation phase evaluations are restricted to forecasts with target end dates (date for which the prediction of instant hospitalizations is being made) within the validation phase (July 27, 2020 to October 31, 2021) - Models passed on: single best stratified by location (states or U.S. national) and model class (regular THieF, SARIMA, THieF ensemble), best regular THieF in last 17 weeks (to account for any models that perform better with more data to train on, especially since they will have a greater amount of data to inform their forecasts during the testing face)

2.7.2 Testing phase

- True time to evaluate the best subset against the two comparison models - All testing phase evaluations are restricted to forecasts with target end dates (date for which the prediction of instant hospitalizations is being made) within the testing phase (November 1, 2021 to October 2, 2022)

CHAPTER 3

RESULTS

3.1 Validation phase

3.2 Testing phase

3.2.1 Summary of comparison models

3.2.2 Overall model performance


	model	wis	mae	cov_50	cov_95	rwis	rmae
1	COVIDhub-4_week_ensemble	1401.30	2085.10	0.63	0.91	0.75	0.83
2	THieF_6wk-4root	1642.50	2358.80	0.40	0.82	0.88	0.94
3	THieF_ensemble-train3	1687.40	2443.40	0.36	0.83	0.90	0.97
4	THieF_ensemble-mean	1700.50	2458.90	0.39	0.84	0.91	0.98
5	THieF_6wk-noTransform	1727.70	2504.80	0.46	0.86	0.92	1.00
6	COVIDhub-baseline	1867.50	2514.10	0.71	0.85	1.00	1.00
7	THieF_12wk-noTransform	2074.10	2908.30	0.37	0.86	1.11	1.16
8	sarima_s7-noTransform	2316.80	3042.40	0.41	0.80	1.24	1.21

Table 3.1. Summary of overall model performance for the US national level, aggregated across forecast week and 1 to 4 week ahead horizons. Prediction interval (PI) coverage rates measure the fraction of times the 50% and 95% PIs covered the true values. A well-calibrated model should achieve close to nominal (0.5 and 0.95, respectively) coverage rates. The “rwis” and “rmae” show relative weighted interval score and relative mean absolute error, which compare each model to the baseline. The baseline is defined to have relative metrics of 1. Models that perform better than the baseline have relative WIS or MAE less than 1. The table is ordered by ascending WIS value. Thus two THieF-4wk models show the best performance out of the THieF models, meeting the baseline for both relative WIS and MAE.

Daily COVID-19 Incident Hospitalizations: observed and forecasted

Selected location(s): United States

Selected forecast date(s): 2021-10-30, 2021-11-27, 2021-12-25, 2022-01-25



ensemble-train3, THieF_ensemble-mean, COVIDhub-baseline, COVIDhub-4_week_ensemble (forecasts)

Figure 3.1.




Figure 3.2. Average WIS by the horizon week for each model for the US national level and the states level. The figures show similar trends in WIS when forecasts are stratified by horizon week compared to being aggregated across horizon weeks. All models show worse performance relative to the baseline at the 1-week ahead horizon, likely a result of being overconfident at this horizon. Once again, the THieF-4wk-4root model performs the best for WIS. Plots for average MAE are not shown due to their similarity to the average WIS plots.




Figure 3.3. Average WIS by the horizon week and pandemic wave for each model for the US national level and the states level. The figures show similar trends in WIS when forecasts are stratified by horizon week and pandemic wave compared to being stratified only across horizon week. Again, all models show worse performance relative to the baseline at the 1-week ahead horizon likely due to being overconfident at this horizon. Model performance is more similar among the top four models during the Alpha wave compared to during the Winter21 wave. Notably, although the THieF-4wk-4root models largely show the lowest WIS, longer horizons during the Winter21 wave saw better performance by the poor-performing THieF-12wk models. Plots for average MAE are not shown due to their similarity to the average WIS plots.




Figure 3.4. Relative WIS by location for each model across all horizons for the entire testing phase. Each square contains the average relative WIS of a model at a particular location, colored based on the relative WIS compare to the baseline. Blue indicates model outperformed the baseline for that location on average, while red indicates the model performed worse than the baseline. Locations are sorted based on cumulative hospitalizations during the testing phase while models are sorted from lowest to highest overall relative WIS, aggregated across states (Table 2). The top three models by relative WIS overall during the testing phase outperformed the baseline in every location.

	model	wis	mae	cov_50	cov_95	rwis	rmae
1	COVIDhub-4_week_ensemble	30.18	44.87	0.64	0.94	0.72	0.79
2	THieF_6wk-4root	34.55	50.12	0.51	0.90	0.83	0.88
3	THieF_ensemble-train3	35.62	51.57	0.56	0.91	0.86	0.91
4	THieF_ensemble-mean	35.91	51.80	0.57	0.91	0.86	0.91
5	THieF_6wk-noTransform	39.51	56.87	0.55	0.88	0.95	1.00
6	COVIDhub-baseline	41.63	56.77	0.79	0.96	1.00	1.00
7	THieF_12wk-noTransform	41.66	60.97	0.51	0.88	1.00	1.07
8	sarima_s7-noTransform	43.54	58.94	0.60	0.87	1.05	1.04

Table 3.2. Summary of overall model performance averaged over the states level. A similar ordering of models is seen, with the THieF-4wk-4root model performing the best and the THieF-12wk models performing the worst. However, the THieF-8wk models perform better in terms of model ranking for the states level. Additionally, relative WIS sees smaller values and 50% and 95% PI coverage see values closer to the nominal levels compared to that at the US national level, although relative MAE is larger with every model performing worse than the baseline.

3.2.3 Model performance by horizon week

3.2.4 Model performance by pandemic wave

3.2.5 Model performance by location

3.2.6 Model performance by forecast week




Figure 3.5. Average h-week ahead WIS and 95% PI coverage for each model for the US national level. Model performance differs from week to week: the WIS is the lowest during weeks with stable and low incident hospitalization values. The THieF-4wk models show the best performance with lower or comparable WIS to the baseline during the majority of the training period with especially good performance for 1-week ahead horizons. Meanwhile, the THieF-12wk showed lower or comparable WIS to the baseline during high incidence and changed points, which have been historically found difficult to forecast accurately. Plots for h-week ahead MAE are not shown due to their similarity to the WIS plots. All THieF models show very variable coverage rates, especially for a 1 week ahead horizon, only the models without a data transformation showing stable coverage starting around February 2021.




Figure 3.6. Average h-week ahead WIS and 95% PI coverage for each model for the states level. Once again, model performance differs from week to week with the lowest WIS values during weeks with stable and low incident hospitalization values and the THieF-4wk models showing the best performance. However, the THieF-12wk did not show good performance during high incidence and changed points for the states level; instead, THieF-12-4root actually showed some strangely large WIS values during April and May 2021. Plots for h-week ahead MAE are not shown due to their similarity to the WIS plots. One-week ahead coverage rates generally improved as time went on, the best models gaining stable and fairly good coverage rates starting in February 2021. For the 4-week ahead horizon, models without a data transformation showed the best performance and stable coverage starting around February 2021, though THieF-4wk-4root and THieF-8wk-4root also saw decent and fairly stable coverage rates after the downward Winter 21 change point.

CHAPTER 4

DISCUSSION

APPENDIX: SUPPLEMENTAL MATERIALS

	model	horizon_wk	wis	mae	cov_50	cov_95	rwis	rmae
1	COVIDhub-4_week_ensemble	1.00	659.80	966.80	0.68	0.94	0.69	0.77
2	THieF_ensemble-train3	1.00	909.60	1348.80	0.37	0.86	0.95	1.07
3	THieF_ensemble-mean	1.00	930.90	1372.90	0.38	0.86	0.98	1.09
4	COVIDhub-baseline	1.00	952.20	1258.80	0.73	0.92	1.00	1.00
5	THieF_6wk-4root	1.00	1007.40	1419.10	0.45	0.79	1.06	1.13
6	THieF_6wk-noTransform	1.00	1028.70	1467.40	0.46	0.88	1.08	1.17
7	sarima_s7-noTransform	1.00	1143.40	1533.90	0.32	0.81	1.20	1.22
8	THieF_12wk-noTransform	1.00	1272.40	1793.60	0.40	0.88	1.34	1.42
9	COVIDhub-4_week_ensemble	2.00	1156.20	1694.70	0.66	0.92	0.72	0.80
10	THieF_6wk-4root	2.00	1413.20	2029.00	0.42	0.85	0.88	0.90
11	THieF_ensemble-train3	2.00	1417.10	2046.90	0.39	0.84	0.88	0.97
12	THieF_ensemble-mean	2.00	1436.90	2070.80	0.40	0.85	0.90	0.98
13	THieF_6wk-noTransform	2.00	1528.50	2177.00	0.47	0.86	0.95	1.03
14	COVIDhub-baseline	2.00	1604.90	2113.30	0.74	0.85	1.00	1.00
15	THieF_12wk-noTransform	2.00	1858.50	2558.30	0.38	0.85	1.16	1.23
16	sarima_s7-noTransform	2.00	1923.10	2536.30	0.41	0.82	1.20	1.20
17	COVIDhub-4_week_ensemble	3.00	1670.60	2528.80	0.60	0.90	0.75	0.84
18	THieF_6wk-4root	3.00	1867.40	2715.30	0.37	0.84	0.84	0.93
19	THieF_ensemble-train3	3.00	1979.60	2845.10	0.35	0.81	0.89	0.93
20	THieF_6wk-noTransform	3.00	1989.00	2883.60	0.46	0.85	0.90	0.90
21	THieF_ensemble-mean	3.00	1989.00	2860.10	0.39	0.83	0.90	0.93
22	COVIDhub-baseline	3.00	2217.00	2993.70	0.71	0.83	1.00	1.00
23	THieF_12wk-noTransform	3.00	2369.90	3313.80	0.37	0.86	1.07	1.13
24	sarima_s7-noTransform	3.00	2738.40	3581.40	0.46	0.80	1.24	1.20
25	COVIDhub-4_week_ensemble	4.00	2173.40	3232.70	0.56	0.89	0.79	0.83
26	THieF_6wk-4root	4.00	2330.00	3341.70	0.36	0.80	0.84	0.88
27	THieF_6wk-noTransform	4.00	2414.90	3567.40	0.43	0.85	0.88	0.94
28	THieF_ensemble-train3	4.00	2501.10	3615.30	0.32	0.80	0.91	0.90
29	THieF_ensemble-mean	4.00	2502.20	3613.60	0.38	0.80	0.91	0.93
30	COVIDhub-baseline	4.00	2761.40	3782.10	0.67	0.80	1.00	1.00
31	THieF_12wk-noTransform	4.00	2852.50	4049.10	0.31	0.86	1.03	1.07
32	sarima_s7-noTransform	4.00	3549.40	4630.40	0.45	0.76	1.28	1.22

Table 1. Summary of model performance stratified by horizon week for the US national level, aggregated across forecast week. The results from this table were used to create Figure 4 that summarized average WIS by horizon week for the US national level and the states level. The table is ordered by ascending horizon week and WIS, with THieF-4wk-4root having the lowest WIS except for the 4 week ahead horizon. The baseline consistently does not have the best WIS.

	model	horizon_wk	wis	mae	cov_50	cov_95	rwis	rmae
1	COVIDhub-4_week_ensemble	1.00	16.70	25.11	0.65	0.96	0.70	0.79
2	THieF_ensemble-train3	1.00	20.32	30.04	0.55	0.93	0.85	0.94
3	THieF_ensemble-mean	1.00	20.76	30.52	0.56	0.93	0.87	0.96
4	THieF_6wk-4root	1.00	22.50	32.97	0.49	0.89	0.94	1.03
5	sarima_s7-noTransform	1.00	23.11	31.92	0.57	0.89	0.97	1.00
6	COVIDhub-baseline	1.00	23.92	31.87	0.80	0.97	1.00	1.00
7	THieF_6wk-noTransform	1.00	25.17	36.41	0.52	0.89	1.05	1.14
8	THieF_12wk-noTransform	1.00	29.17	42.20	0.49	0.89	1.22	1.32
9	COVIDhub-4_week_ensemble	2.00	25.42	37.90	0.66	0.95	0.69	0.77
10	THieF_6wk-4root	2.00	30.10	43.94	0.52	0.91	0.81	0.89
11	THieF_ensemble-train3	2.00	30.66	44.49	0.56	0.91	0.83	0.90
12	THieF_ensemble-mean	2.00	31.04	44.80	0.58	0.91	0.84	0.90
13	THieF_6wk-noTransform	2.00	34.99	50.27	0.56	0.88	0.95	1.02
14	COVIDhub-baseline	2.00	36.95	49.48	0.80	0.96	1.00	1.00
15	sarima_s7-noTransform	2.00	37.32	50.28	0.60	0.87	1.01	1.02
16	THieF_12wk-noTransform	2.00	37.35	54.31	0.53	0.88	1.01	1.10
17	COVIDhub-4_week_ensemble	3.00	35.22	52.52	0.64	0.94	0.73	0.79
18	THieF_6wk-4root	3.00	38.98	56.36	0.52	0.90	0.80	0.85
19	THieF_ensemble-train3	3.00	41.60	59.88	0.56	0.90	0.86	0.90
20	THieF_ensemble-mean	3.00	41.85	60.07	0.58	0.90	0.86	0.91
21	THieF_6wk-noTransform	3.00	45.11	64.83	0.56	0.87	0.93	0.98
22	THieF_12wk-noTransform	3.00	46.50	68.16	0.52	0.87	0.96	1.03
23	COVIDhub-baseline	3.00	48.54	66.30	0.79	0.95	1.00	1.00
24	sarima_s7-noTransform	3.00	51.31	69.07	0.60	0.86	1.06	1.04
25	COVIDhub-4_week_ensemble	4.00	44.39	65.43	0.63	0.93	0.76	0.80
26	THieF_6wk-4root	4.00	47.53	68.48	0.52	0.88	0.81	0.84
27	THieF_ensemble-train3	4.00	51.04	73.47	0.55	0.89	0.88	0.90
28	THieF_ensemble-mean	4.00	51.11	73.39	0.56	0.88	0.88	0.90
29	THieF_6wk-noTransform	4.00	53.80	77.46	0.55	0.86	0.92	0.95
30	THieF_12wk-noTransform	4.00	54.52	80.61	0.50	0.87	0.93	0.99
31	COVIDhub-baseline	4.00	58.35	81.23	0.77	0.95	1.00	1.00
32	sarima_s7-noTransform	4.00	63.89	86.44	0.60	0.85	1.09	1.06

Table 2. Summary of model performance stratified by horizon week for the states level, aggregated across forecast week. The results from this table were used to create Figure 4 that summarized average WIS by horizon week for the US national level and the states level. The table is ordered by ascending horizon week and WIS, with THieF-4wk-4root having the lowest WIS except for the 3 week (although it is second lowest) and 4 week ahead horizon. The baseline again consistently does not have the best WIS, performing progressively worse compared to the other models for each additional week ahead horizon.

	model	horizon_wk	wave	wis	mae	cov_50	cov_95
1	COVIDhub-4_week_ensemble	1.00	ba4_ba5	277.80	378.70	0.89	1.00
2	THieF_6wk-4root	1.00	ba4_ba5	350.40	542.90	0.60	0.99
3	THieF_ensemble-train3	1.00	ba4_ba5	399.10	685.30	0.46	1.00
4	THieF_ensemble-mean	1.00	ba4_ba5	433.60	740.30	0.49	1.00
5	COVIDhub-baseline	1.00	ba4_ba5	450.80	486.90	0.92	1.00
6	THieF_6wk-noTransform	1.00	ba4_ba5	516.50	822.50	0.60	1.00
7	sarima_s7-noTransform	1.00	ba4_ba5	517.80	828.60	0.41	0.99
8	THieF_12wk-noTransform	1.00	ba4_ba5	790.70	1319.20	0.35	1.00
9	COVIDhub-4_week_ensemble	2.00	ba4_ba5	411.40	572.50	0.85	1.00
10	THieF_6wk-4root	2.00	ba4_ba5	518.80	813.00	0.61	1.00
11	THieF_ensemble-train3	2.00	ba4_ba5	578.60	966.60	0.49	1.00
12	THieF_ensemble-mean	2.00	ba4_ba5	621.80	1030.30	0.50	1.00
13	COVIDhub-baseline	2.00	ba4_ba5	654.20	741.00	0.98	1.00
14	THieF_6wk-noTransform	2.00	ba4_ba5	727.80	1145.90	0.59	1.00
15	sarima_s7-noTransform	2.00	ba4_ba5	729.70	1169.60	0.53	1.00
16	THieF_12wk-noTransform	2.00	ba4_ba5	1012.60	1677.50	0.38	1.00
17	COVIDhub-4_week_ensemble	3.00	ba4_ba5	563.30	824.70	0.77	1.00
18	THieF_6wk-4root	3.00	ba4_ba5	683.90	1069.10	0.59	1.00
19	THieF_ensemble-train3	3.00	ba4_ba5	784.00	1311.50	0.45	1.00
20	COVIDhub-baseline	3.00	ba4_ba5	825.00	1038.90	1.00	1.00
21	THieF_ensemble-mean	3.00	ba4_ba5	825.00	1371.30	0.50	1.00
22	THieF_6wk-noTransform	3.00	ba4_ba5	911.00	1464.40	0.59	1.00
23	sarima_s7-noTransform	3.00	ba4_ba5	953.50	1514.20	0.57	1.00
24	THieF_12wk-noTransform	3.00	ba4_ba5	1248.70	2097.10	0.38	1.00
25	COVIDhub-4_week_ensemble	4.00	ba4_ba5	698.90	1034.30	0.77	1.00
26	THieF_6wk-4root	4.00	ba4_ba5	801.90	1210.30	0.64	1.00
27	THieF_ensemble-train3	4.00	ba4_ba5	969.80	1585.10	0.48	1.00
28	COVIDhub-baseline	4.00	ba4_ba5	972.00	1320.50	1.00	1.00
29	THieF_ensemble-mean	4.00	ba4_ba5	1010.80	1651.70	0.53	1.00
30	THieF_6wk-noTransform	4.00	ba4_ba5	1034.00	1631.70	0.61	1.00
31	sarima_s7-noTransform	4.00	ba4_ba5	1169.60	1821.80	0.60	1.00
32	THieF_12wk-noTransform	4.00	ba4_ba5	1501.80	2581.80	0.34	1.00
33	THieF_ensemble-mean	1.00	delta	239.10	305.40	0.86	1.00
34	THieF_ensemble-train3	1.00	delta	247.90	326.80	0.86	1.00
35	sarima_s7-noTransform	1.00	delta	270.90	454.30	0.29	1.00
36	THieF_6wk-4root	1.00	delta	327.60	536.40	0.71	1.00
37	THieF_6wk-noTransform	1.00	delta	332.50	563.70	0.43	1.00
38	THieF_12wk-noTransform	1.00	delta	381.80	523.00	1.00	1.00
39	THieF_6wk-noTransform	2.00	delta	286.50	260.10	1.00	1.00
40	sarima_s7-noTransform	2.00	delta	290.70	240.00	1.00	1.00
41	THieF_ensemble-mean	2.00	delta	313.60	215.30	1.00	1.00
42	THieF_ensemble-train3	2.00	delta	316.20	207.60	1.00	1.00
43	THieF_6wk-4root	2.00	delta	324.40	280.40	1.00	1.00
44	THieF_12wk-noTransform	2.00	delta	487.90	404.70	1.00	1.00
45	THieF_6wk-noTransform	3.00	delta	396.00	511.10	0.86	1.00
46	sarima_s7-noTransform	3.00	delta	486.30	701.80	1.00	1.00
47	THieF_ensemble-mean	3.00	delta	509.60	800.20	1.00	1.00
48	THieF_ensemble-train3	3.00	delta	509.60	797.20	1.00	1.00

	model	horizon_wk	wave	wis	mae	cov_50	cov_95	rw
1	COVIDhub-4_week_ensemble	1.00	ba4_ba5	8.91	13.43	0.72	0.99	0.6
2	THieF_6wk-4root	1.00	ba4_ba5	9.95	14.83	0.62	0.97	0.6
3	THieF_ensemble-train3	1.00	ba4_ba5	10.48	15.62	0.67	0.99	0.7
4	THieF_ensemble-mean	1.00	ba4_ba5	11.10	16.47	0.69	0.99	0.7
5	sarima_s7-noTransform	1.00	ba4_ba5	12.81	19.00	0.68	0.98	0.8
6	THieF_6wk-noTransform	1.00	ba4_ba5	13.93	20.54	0.70	0.99	0.9
7	COVIDhub-baseline	1.00	ba4_ba5	14.49	16.59	0.92	1.00	1.0
8	THieF_12wk-noTransform	1.00	ba4_ba5	18.37	28.32	0.61	0.99	1.2
9	COVIDhub-4_week_ensemble	2.00	ba4_ba5	11.85	17.64	0.76	1.00	0.6
10	THieF_6wk-4root	2.00	ba4_ba5	13.28	19.07	0.69	0.99	0.6
11	THieF_ensemble-train3	2.00	ba4_ba5	14.15	20.34	0.74	0.99	0.7
12	THieF_ensemble-mean	2.00	ba4_ba5	14.86	21.20	0.75	0.99	0.7
13	sarima_s7-noTransform	2.00	ba4_ba5	17.46	25.07	0.75	0.99	0.8
14	THieF_6wk-noTransform	2.00	ba4_ba5	17.92	25.69	0.77	1.00	0.9
15	COVIDhub-baseline	2.00	ba4_ba5	19.69	21.67	0.97	1.00	1.0
16	THieF_12wk-noTransform	2.00	ba4_ba5	21.52	32.82	0.68	1.00	1.0
17	COVIDhub-4_week_ensemble	3.00	ba4_ba5	15.22	22.86	0.76	0.99	0.6
18	THieF_6wk-4root	3.00	ba4_ba5	16.82	23.30	0.72	0.99	0.7
19	THieF_ensemble-train3	3.00	ba4_ba5	17.95	25.29	0.76	1.00	0.7
20	THieF_ensemble-mean	3.00	ba4_ba5	18.68	26.22	0.78	1.00	0.7
21	THieF_6wk-noTransform	3.00	ba4_ba5	21.59	31.20	0.78	1.00	0.9
22	sarima_s7-noTransform	3.00	ba4_ba5	22.06	31.34	0.78	0.99	0.9
23	COVIDhub-baseline	3.00	ba4_ba5	23.93	27.49	0.97	1.00	1.0
24	THieF_12wk-noTransform	3.00	ba4_ba5	25.38	39.28	0.70	1.00	1.0
25	COVIDhub-4_week_ensemble	4.00	ba4_ba5	18.26	27.51	0.76	0.99	0.6
26	THieF_6wk-4root	4.00	ba4_ba5	19.92	26.57	0.75	1.00	0.7
27	THieF_ensemble-train3	4.00	ba4_ba5	21.50	29.63	0.78	1.00	0.7
28	THieF_ensemble-mean	4.00	ba4_ba5	22.13	30.41	0.79	1.00	0.8
29	THieF_6wk-noTransform	4.00	ba4_ba5	24.17	34.33	0.80	1.00	0.8
30	sarima_s7-noTransform	4.00	ba4_ba5	25.95	36.72	0.80	0.99	0.9
31	COVIDhub-baseline	4.00	ba4_ba5	27.46	32.80	0.97	1.00	1.0
32	THieF_12wk-noTransform	4.00	ba4_ba5	29.06	45.41	0.69	1.00	1.0
33	sarima_s7-noTransform	1.00	delta	10.38	14.69	0.70	0.97	0.3
34	THieF_ensemble-mean	1.00	delta	11.60	18.76	0.64	0.99	0.3
35	THieF_ensemble-train3	1.00	delta	11.87	19.30	0.60	0.99	0.3
36	THieF_6wk-4root	1.00	delta	14.16	22.68	0.51	0.95	0.4
37	THieF_12wk-noTransform	1.00	delta	14.89	23.52	0.52	0.95	0.4
38	THieF_6wk-noTransform	1.00	delta	15.57	25.83	0.51	0.97	0.4
39	sarima_s7-noTransform	2.00	delta	14.64	20.80	0.69	0.98	0.2
40	THieF_ensemble-mean	2.00	delta	16.33	26.27	0.65	0.99	0.3
41	THieF_ensemble-train3	2.00	delta	16.52	26.76	0.64	0.98	0.3
42	THieF_6wk-4root	2.00	delta	17.99	29.88	0.54	0.98	0.3
43	THieF_12wk-noTransform	2.00	delta	19.35	30.10	0.60	0.95	0.3
44	THieF_6wk-noTransform	2.00	delta	19.58	32.20	0.62	0.97	0.3
45	sarima_s7-noTransform	3.00	delta	23.20	35.71	0.60	0.97	0.3
46	THieF_ensemble-mean	23 3.00	delta	25.61	41.49	0.58	0.97	0.3
47	THieF_ensemble-train3	3.00	delta	25.77	41.94	0.57	0.98	0.3
48	THieF_6wk-4root	3.00	delta	27.03	44.49	0.52	0.99	0.3

BIBLIOGRAPHY

- [1] COVID-19 Forecast Hub. Ensemble model. <https://covid19forecasthub.org/doc/ensemble/>. Accessed: 2023-03-15.
- [2] Cramer, Estee Y, Huang, Yuxin, Wang, Yijin, Ray, Evan L, Cornell, Matthew, Bracher, Johannes, Brennen, Andrea, Castro Rivadeneira, Alvaro J, Gerding, Aaron, House, Katie, Jayawardena, Dasuni, Kanji, Abdul H, Khandelwal, Ayush, Le, Khoa, Niemi, Jarad, Stark, Ariane, Shah, Apurv, Wattanachit, Nutch, Zorn, Martha W, Reich, Nicholas G, and Consortium, US COVID-19 Forecast Hub. The united states covid-19 forecast hub dataset. *medRxiv* (2021).
- [3] Cramer, Estee Y., Ray, Evan L., Lopez, Velma K., Bracher, Johannes, Brennen, Andrea, Rivadeneira, Alvaro J. Castro, Gerding, Aaron, Gneiting, Tilmann, Katie H. House, Yuxin Huang, Dasuni Jayawardena, Kanji, Abdul H., Khandelwal, Ayush, Le, Khoa, Mühlemann, Anja, Niemi, Jarad, Shah, Apurv, Stark, Ariane, Wang, Yijin, Wattanachit, Nutch, ..., and Reich, Nicholas G. Evaluation of individual and ensemble probabilistic forecasts of covid-19 mortality in the us. *Proceedings of the National Academy of Sciences* 119, 15 (April 2022), e2113561119.
- [4] Donovan, Doug. Johns hopkins covid-19 data hub ends after three years. *The Hub* (March 2023).
- [5] Gibson, Graham C., Reich, Nicholas G., and Sheldon, Daniel. Real-time mechanistic bayesian forecasts of covid-19 mortality. *medRxiv : the preprint server for health sciences* (December 2020).
- [6] Reich, Nicholas G., Tibshirani, Ryan J., Ray, Evan L., and Rosenfeld, Roni. On the predictability of covid-19. <https://forecasters.org/blog/2021/09/28/on-the-predictability-of-covid-19/>, September 2021. Accessed: 2022-04-06.
- [7] U.S. Department of Health and Human Services. Covid-19 reported patient impact and hospital capacity by state timeseries. <https://healthdata.gov/Hospital/COVID-19-Reported-Patient-Impact-and-Hospital-Capa/g62h-syeh>. Accessed: 2023-03-15.